

# 平行程式設計 hw3 report

資工四 B10502136 邱垣盛

## 1 Implementation

### 1.1 CUDA Implementation

The program's implementation revolves around migrating the per-pixel rendering task, originally handled by nested `for` loops on the CPU, to a massively parallel CUDA kernel named `renderKernel_tiled`. All core computational logic, including the Mandelbulb distance function (`md`), scene mapping (`map`), ray tracing (`trace`), and shading functions (`calcNor`, `softshadow`), were converted into `__device__` functions. The host-side `main` function is responsible for setup and teardown: it allocates device memory for the output image (`d_image`), copies the final result back to the host (`h_image`), and manages the kernel launch.

A key part of the implementation is the efficient handling of global scene data. Instead of passing many parameters to the kernel, scene-wide, read-only values like the camera position (`d_camera_pos`), target (`d_target_pos`), resolution (`d_iResolution`), and fractal parameters (`d_power`, `d_md_iter`) are copied from the host to `__constant__` device memory before the kernel launch. This data is then broadcast to all threads with low latency via a dedicated cache, which is much faster than repeated global memory reads. The kernel itself is launched with a grid of blocks, and its primary arguments are the pointer to the device image buffer (`d_image`) and the image dimensions.

### 1.2 Task Partitioning

Task partitioning is achieved using a dynamic load-balancing model built on spatial tiles. The image is divided into `16x16` pixel tiles. The kernel is launched with a 1D grid of thread blocks, where each block contains `16x16` (256) threads. A global `unsigned int` in device memory, `g_tileCounter`, acts as a work queue. Inside a `while(true)` loop, the first thread in each block (`threadIdx.x == 0 && threadIdx.y == 0`) atomically increments this counter to claim a unique `tileId`. All threads in that block then use this `tileId` to calculate the corresponding pixel coordinates they are responsible for rendering.

This dynamic task stealing model is highly effective for this workload. When a block finishes rendering an easy tile (like empty sky), it simply loops and claims the next available `tileId` from the counter. This prevents blocks from becoming idle, maximizing GPU utilization. The loop terminates when a block claims a `tileId` that is greater than or equal to the `totalTiles`.

This exit condition is checked by the first thread and broadcast to the rest of the block using `__shared__` memory (`shouldQuit`) and synchronization barriers (`__syncthreads()`). This approach is superior to a static 2D grid, one block per tile, where workload imbalance would cause the entire kernel to wait for the single block rendering the most complex part of the fractal.

### 1.3 Optimization Skills

Occupancy tuning launches much more blocks than the number of streaming multiprocessors available. This ensures the GPU always has work available when blocks stall on memory operations, hiding latency and maintaining high throughput.

Shared memory reduces atomic contention by having only one thread per block perform atomic operations and broadcast results to all threads. This reduces atomic memory operations by a factor of 256 compared to having every thread access the global counter independently.

The implementation reduces precision from double (64-bit) to float (32-bit) throughout, as Mandelbulb rendering doesn't require double precision. Float operations execute approximately twice as fast on GPU architectures and consume half the memory bandwidth. Memory is preinitialized with `cudaMemset()` to ensure clean output without requiring conditional writes in the kernel.

The packed RGBA write pattern ensures threads in a warp access consecutive memory addresses, maximizing memory bandwidth utilization to up to 128 bytes per transaction. This coalesced access pattern is critical for achieving high memory throughput on GPU architectures.

## 2 Analysis

### 2.1 nvprof


Based on the profiling results, the optimal GPU kernel execution time was 10650.00 ms, achieved with the 8x8 block size configuration in test case 8.

```
==1192068== Profiling application: ./hw3 -1.2 -0.51 -0.8 -0.271 -0.299 -0.379 4096 4096 08.png
==1192068== Profiling result:
   Type Time(%) Time Calls Avg Min Max Name
GPU activities: 99.77% 9.23567s 1 9.23567s 9.23567s 9.23567s renderKernel_tiled(unsigned char*, int, int)
0.23% 21.263ms 1 21.263ms 21.263ms 21.263ms [CUDA memcpy DtoH]
0.00% 81.759us 1 81.759us 81.759us 81.759us [CUDA memset]
0.00% 4.2240us 4 1.0560us 1.0560us 1.0560us [CUDA memcpy HtoD]
API calls: 98.65% 9.23567s 1 9.23567s 9.23567s 9.23567s cudaDeviceSynchronize
1.10% 103.21ms 4 25.803ms 5.4540us 103.11ms cudaMemcpyToSymbol
0.23% 21.907ms 1 21.907ms 21.907ms 21.907ms cudaMemcpy
0.00% 376.42us 114 3.3010us 71ns 176.72us cuDeviceGetAttribute
0.00% 360.01us 1 360.01us 360.01us 360.01us cudaFree
0.00% 184.35us 1 184.35us 184.35us 184.35us cudaMalloc
0.00% 56.953us 1 56.953us 56.953us 56.953us cudaMemcpy
0.00% 37.226us 2 18.613us 5.9950us 31.231us cudaEventRecord
0.00% 36.706us 1 36.706us 36.706us 36.706us cudaLaunchKernel
0.00% 32.588us 2 16.294us 534ns 32.054us cudaEventCreate
0.00% 24.446us 1 24.446us 24.446us 24.446us cudaEventElapsedTime
0.00% 9.3300us 1 9.3300us 9.3300us 9.3300us cuDeviceGetName
0.00% 9.2150us 2 4.6070us 578ns 8.6370us cudaEventDestroy
0.00% 7.5630us 1 7.5630us 7.5630us 7.5630us cuDeviceGetAttribute
0.00% 7.2280us 1 7.2280us 7.2280us 7.2280us cuDeviceGetPCIBusId
0.00% 6.9660us 1 6.9660us 6.9660us 6.9660us cudaGetDevice
0.00% 1.4080us 3 469ns 97ns 1.2080us cuDeviceGetCount
0.00% 530ns 2 265ns 131ns 399ns cuDeviceGet
0.00% 295ns 1 295ns 295ns 295ns cuDeviceTotalMem
0.00% 227ns 1 227ns 227ns 227ns cuDeviceGetUuid
0.00% 193ns 1 193ns 193ns 193ns cuModuleGetLoadingMode
0.00% 163ns 1 163ns 163ns 163ns cudaGetLastError
Rendering with 960 blocks of 16x16 threads (Dynamic Tiling)
Total tiles: 256 x 256 = 65536
Kernel time: 9235.833 ms
Throughput: 1.817 Mpix/s
```

## 2.2 ncu

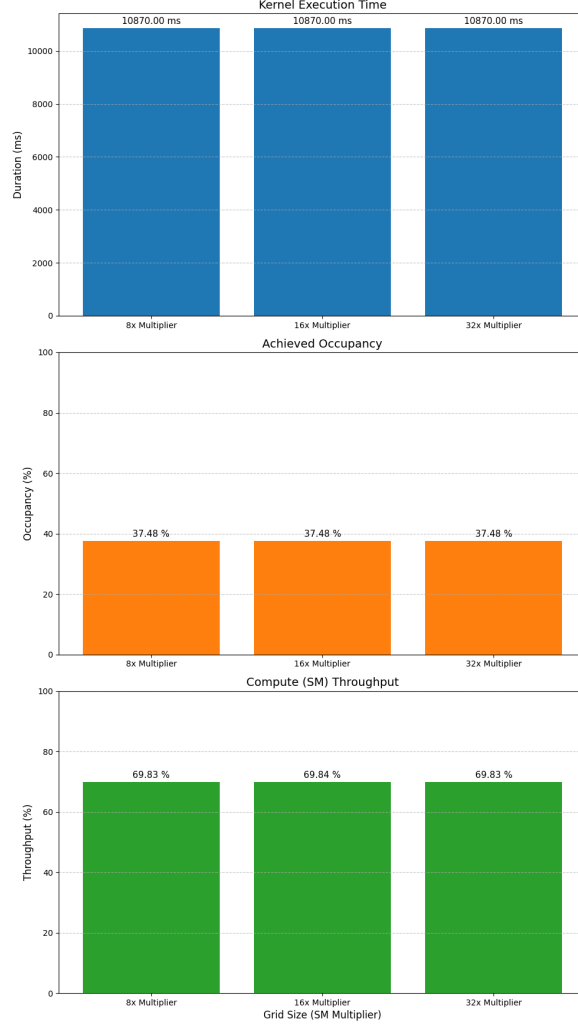
**Different block size** Based on the profiling results, the optimal kernel execution time was 10650.00 ms, achieved with an 8x8 block size.

The Nsight Compute plots show this performance difference clearly. The 4x4 block was inefficient, with low Achieved Occupancy (18.74%). Switching to 8x8 doubled occupancy to 37.5% and boosted SM Throughput from 56.68% to 69.58%, resulting in the fastest execution time. The 16x16 block offered no significant further improvement.



**Different grid size** Based on profiling, the kernel execution time was 10870.00 ms.

The Nsight Compute plots show that performance saturated with the 8x multiplier. Increasing the grid size to 16x or 32x provided no change in execution time, occupancy (37.48%), or SM throughput (69.8%), as the GPU was already fully utilized.



## 2.3 Additions

The dynamic load-balancing strategy is highly effective. The grid size analysis proves that performance saturates at the 8x multiplier; increasing the grid size further has no impact. This design successfully manages the unbalanced workload and removes the need for precise grid size tuning.

This workload is clearly compute-bound, not memory-bound, as SM Throughput reached 70% while DRAM Throughput was near 0%. The 8x8 block size was optimal, providing the fastest execution time (10650.00 ms) by doubling the achieved occupancy (37.5%) compared to the inefficient 4x4 block.

## 3 Conclusion

### 3.1 Difficulty

The primary difficulty encountered was managing the highly unbalanced workload inherent in rendering a fractal. The computational cost varies dramatically per pixel; regions of empty space

are computed quickly, while pixels near the complex fractal boundary require extensive iteration. A simple, static partitioning of the image among thread blocks would result in severe inefficiency, with most blocks finishing early and sitting idle while a few blocks handled the complex areas. This was overcome by implementing a dynamic tiling system. Using a global atomic counter (`g_tileCounter`), thread blocks could "steal" the next available tile of work, ensuring that GPU resources remained highly utilized throughout the rendering process.

A secondary challenge involved the low-level performance optimization of the core distance-estimation function (`md`) on the GPU. Porting the mathematical operations from the CPU required careful optimization to maximize throughput. This involved replacing expensive `pow` operations with faster, explicit multiplications (e.g., `r2`, `r4`, `r8`) and using the `sincosf` intrinsic to calculate sine and cosine in a single instruction. Additionally, optimizing memory access patterns, such as using `__constant__` memory for global parameters and packing final pixel data into a `uint32_t` for coalesced writes, was essential to prevent memory bottlenecks.

## 3.2 Feedback

This assignment was my first experience with CUDA, and my key takeaway is that GPU acceleration is not automatic. I learned that effective parallel programming requires a deep understanding of the hardware's constraints, such as occupancy, register limits, and shared memory, rather than just moving code to the GPU. This fractal's unbalanced workload also taught me the importance of implementing a dynamic load-balancing strategy, like the tiling system used. Finally, I realized that even after a correct parallel architecture is designed, performance is not guaranteed. It requires significant fine-tuning of launch parameters, such as the block size and grid number, to truly maximize throughput.